

Homework #3, EE556, Fall 2018

Due: 9/27/18; hand in Problems 1, 2, and 6

Problem #1

A classifier uses four linear discriminant functions in the plane: $g_1 = x_1$, $g_2 = x_2$, $g_3 = x_2 + x_1 - 3$, and $g_4 = x_2 - x_1 + 1$.

- i) Sketch the line boundaries in pattern space and label each region with a 4-bit binary codeword.
- ii) Make a logical table, identifying for each codeword whether or not there is an associated region.
- iii) Sketch the **four**-dimensional hypercube in state space. (**Don't panic !** Use two 3-dimensional projections, one yielding a 3D cube for the case $T_1 = 0$ and the other giving a 3D cube for $T_1 = 1$). Label each vertex with its corresponding binary codeword. Draw a curved line joining each vertex in the first cube with the corresponding vertex in the second cube.
- ii) Consider the discriminant function $y = \text{sgn}(\sum_{i=1}^4 T_i - 2.5)$. Sketch the decision region induced by this rule in the (2D) feature space.

Problem #2

Construct a multilayer perceptron that solves the N -bit parity problem. (Recall this problem from homework 2).

Problem #3

- i) Consider an MLP with I inputs, J hidden units, and K output units (a single hidden layer). What is the space complexity of the network ? (Include the the storage required for MLP parameters, training data, and and any additional storage needed during training).
- ii) What is the computational complexity of backpropagation training in batch gradient descent mode ?

Problem #4

Consider a standard multilayer perceptron. Show that if the sign on every weight is flipped the operation of the network remains unchanged (a type of “polar symmetry”) – does this property

extend to changing the signs of the inputs to the network ?

Problem #5

Derive the learning rule for updating an input-to-hidden unit weight for a single-hidden layer MLP that uses a “softmax function” in the output layer and the cross entropy criterion for training (both “softmax” and “cross entropy” are discussed in lecture).

Problem #6

Consider the problem of neural network *inversion*, wherein, given a fixed network and target output values, the objective is to learn the associated *input* patterns which, when forward-propagated through the network, produce outputs that well-approximate the targets. Sketch a method (an optimization technique ?) that approximately achieves NN inversion. Also suggest some possible applications for NN inversion.