

EE 556

HW #4 Solutions

#2)

(6.6, Haykin)

$K = [K(\underline{x}_i, \underline{x}_j)]$ is a square matrix.

Therefore, it can be written as:

$K = Q \Lambda Q^T$, where Λ is a diagonal matrix w/ the eigenvalues of K and Q is an orthogonal matrix whose columns are the associated eigenvectors (this is often called a "similarity transform").
Because $K(\underline{x}_i, \underline{x}_j)$ is positive definite, all the eigenvalues are non-negative.

We can write:

$$K(\underline{x}_i, \underline{x}_j) = (Q \Lambda Q^T)_{ij} = \sum_{l=1}^m (Q)_{il} \Lambda_{ll} (Q^T)_{lj}$$
$$= \sum_{l=1}^m Q_{il} \Lambda_{ll} Q_{jl} \quad (\text{since for an orthogonal matrix, } Q = Q^T)$$

Let \underline{u}_i denote the i -th row of the matrix Q . (Note that \underline{u}_i is NOT an eigenvector of $K \dots$). (The columns \underline{e} are the eigenvectors...).

$$\text{Then, } K(\underline{x}_i, \underline{x}_j) = \underline{u}_i^T \Lambda \underline{u}_j$$
$$= (\Lambda^{1/2} \underline{u}_i)^T (\Lambda^{1/2} \underline{u}_j)$$

By definition, $K(\underline{x}_i, \underline{x}_i) = \underline{\phi}^T(\underline{x}_i) \underline{\phi}(\underline{x}_i)$.

Therefore, we have: $\underline{\phi}(\underline{x}_i) = \Lambda^{1/2} \underline{u}_i$, i.e.
the mapping from the input space to the feature
space for a kernel SVM is given by:

$$\phi: \underline{x}_i \rightarrow \Lambda^{1/2} \underline{u}_i \quad \left(\begin{array}{l} \text{the mapping for the} \\ \text{training vectors} \dots \end{array} \right)$$

#3)

We employ proof by contradiction. Suppose there were two *distinct* minimum length solution vectors a_1 and a_2 with $a_1^t y > 0$ and $a_2^t y > 0$. Then necessarily we would have $\|a_1\| = \|a_2\|$ (otherwise the longer of the two vectors would not be a *minimum* length solution). Next consider the average vector $a_o = \frac{1}{2}(a_1 + a_2)$. We note that

$$a_o^t y = \frac{1}{2}(a_1 + a_2)^t y = \frac{1}{2}a_1^t y + \frac{1}{2}a_2^t y \geq 0,$$

and thus a_o is indeed a solution vector. Its length is

$$\|a_o\| = \|1/2(a_1 + a_2)\| = 1/2\|a_1 + a_2\| \leq 1/2(\|a_1\| + \|a_2\|) = \|a_1\| = \|a_2\|,$$

where we used the triangle inequality for the Euclidean metric. Thus a_o is a solution vector such that $\|a_o\| \leq \|a_1\| = \|a_2\|$. But by our hypothesis, a_1 and a_2 are minimum length solution vectors. Thus we must have $\|a_o\| = \|a_1\| = \|a_2\|$, and thus

$$\frac{1}{2}\|a_1 + a_2\| = \|a_1\| = \|a_2\|.$$

We square both sides of this equation and find

$$\frac{1}{4}\|a_1 + a_2\|^2 = \|a_1\|^2$$

or

$$\frac{1}{4}(\|a_1\|^2 + \|a_2\|^2 + 2a_1^t a_2) = \|a_1\|^2.$$

We regroup and find

$$\begin{aligned} 0 &= \|a_1\|^2 + \|a_2\|^2 - 2a_1^t a_2 \\ &= \|a_1 - a_2\|^2, \end{aligned}$$

and thus $a_1 = a_2$, contradicting our hypothesis. Therefore, the minimum-length solution vector is unique.

#4) Recall Wolfe-Dual problem:

$$\max_{\underline{\lambda}} \left\{ -\frac{1}{2} \sum_{x, x'} \lambda_x \lambda_{x'} t_x t_{x'} \underline{x}^T \underline{x} + \sum_x \lambda_x \right\}$$

$$\text{s.t. } \sum_x \lambda_x t_x = 0, \quad \underline{\lambda} \geq 0$$

If the data vectors are orthogonal, $\underline{x}^T \underline{x} = \delta_{x'x} \Rightarrow$
 $\begin{matrix} \text{unit norm} \\ \text{orthonormal} \end{matrix}$

the function to be maximized reduces to:

$$-\frac{1}{2} \sum_x \lambda_x^2 + \sum_x \lambda_x$$

For now, ignore $\underline{\lambda} \geq 0$ & take the equality constraint into account via a Lagrange multiplier, μ .

I.e., form

$$L = -\frac{1}{2} \sum_x \lambda_x^2 + \sum_x \lambda_x - \mu \left(\sum_x \lambda_x t_x \right)$$

$$\frac{\partial L}{\partial \lambda_z} = -\lambda_z + 1 - \mu t_z = 0 \Rightarrow$$

$$\lambda_z = 1 - \mu t_z$$

Plugging into $\sum_x \lambda_x t_x = 0$ gives

$$\sum_x (1 - \mu t_x) t_x \stackrel{=0}{=} \Rightarrow \mu = \frac{\sum_x t_x}{N}, \quad N$$

the # of data points.

Since $t_x \in \{-1, +1\}$, $-1 < \mu < 1 \Rightarrow \lambda_x \geq 0, \forall x$,
 i.e. inequality constraint is automatically satisfied.

This also means that all training points are support vectors, in this case.

The solution thus has the form

$$\underline{w} = \sum_x \lambda_x t_x \underline{x} = \sum_x \left(1 - t_x \left(\frac{\sum_{x'} t_{x'}}{N}\right)\right) t_x \underline{x}$$

$$+ \\ w_0 = \cancel{t_x} - \underline{w}^T \underline{x}, \text{ any } x.$$

$$= t_x - \lambda_x t_x = t_x - (1 - \mu t_x) t_x$$

$$= \mu = \boxed{\frac{\sum_x t_x}{N}}$$

Also, we can rewrite \underline{w} as:

$$\underline{w} = \sum_x (1 - \mu t_x) t_x \underline{x}$$

$$= \boxed{\sum_x (t_x - \mu) \underline{x}}$$

This is related to Hebbian learning...

#5) Theorem: Given M linearly independent vectors $\underline{x}_1, \dots, \underline{x}_M$, $\underline{x}_i \in \mathbb{R}^N$, $M < N$,
 \exists a vector \underline{w} s.t. $\underline{w}^T \underline{x}_i > 0$, $i=1, \dots, M$

Proof: The theorem statement is equivalent to the statement that

$\underline{X}^T \underline{w} = \underline{b} > \underline{0}$, that is, the right-hand side is a vector with strictly positive entries.

Here, $\underline{X} = \begin{bmatrix} \underline{x}_1 & \dots & \underline{x}_M \end{bmatrix}$.
 $N \times M$

The matrix \underline{X} has full column rank, M .

Moreover, its row rank is also M . (see, e.g. [Strang]).

This means that \underline{X}^T has an M -dimensional column basis. However, note that the columns of \underline{X}^T are M -dimensional vectors. This implies that the columns of \underline{X}^T span $\mathbb{R}^M \Rightarrow$ any vector $\underline{b} \in \mathbb{R}^M$ is in the ^{column} ~~row~~ space of \underline{X}^T , including \underline{b} with all positive entries.

Thus, $\exists \underline{w}$ s.t. $\underline{X}^T \underline{w} = \underline{b} > \underline{0}$.

```

% XOR input for x1 and x2
input = [0 0; 0 1; 1 0; 1 1];
% Desired output of XOR
output = [0; 1; 1; 0];
% Initialize the bias
bias = [-1 -1 -1];
% Learning coefficient
coeff = 0.7;
% Number of learning iterations
iterations = 10000;
% Calculate weights randomly using seed.
rand('state',sum(100*clock));
weights = -1 + 2.*rand(3,3);

```

Matlab
code for
XOR problem --
updates weights with
each sample presentation

```

for i = 1:iterations
    out = zeros(4,1);
    numIn = length(input(:,1));
    for j = 1:numIn
        % Hidden layer
        H1 = bias(1,1)*weights(1,1) + input(j,1)*weights(1,2) + input(j,2)*weights(1,3);

        % Send data through sigmoid function  $1/(1+e^{-x})$ 
        % Note that sigma is a different m file
        % that I created to run this operation
        x2(1) = 1/(1 + exp(-H1));
        H2 = bias(1,2)*weights(2,1) + input(j,1)*weights(2,2) + input(j,2)*weights(2,3);
        x2(2) = 1/(1+exp(-H2));


        % Output layer
        x3_1 = bias(1,3)*weights(3,1) + x2(1)*weights(3,2) + x2(2)*weights(3,3);
        out(j) = 1/(1+exp(-x3_1));

        % Adjust delta values of weights
        % For output layer:
        %  $\delta(w_i) = x_i * \delta$ 
        %  $\delta = (1 - \text{actual output}) * (\text{desired output} - \text{actual output})$ 
        delta3_1 = out(j)*(1-out(j))*(output(j)-out(j));
        % Propagate the delta backwards into hidden layers
        delta2_1 = x2(1)*(1-x2(1))*weights(3,2)*delta3_1;
        delta2_2 = x2(2)*(1-x2(2))*weights(3,3)*delta3_1;


        % Add weight changes to original weights
        % And use the new weights to repeat process.
        %  $\delta \text{ weight} = \text{coeff} * x * \delta$ 
        for k = 1:3
            if k == 1 % Bias cases
                weights(1,k) = weights(1,k) + coeff*bias(1,1)*delta2_1;
                weights(2,k) = weights(2,k) + coeff*bias(1,2)*delta2_2;
                weights(3,k) = weights(3,k) + coeff*bias(1,3)*delta3_1;
            else % When k=2 or 3 input cases to neurons
                weights(1,k) = weights(1,k) + coeff*input(j,k-1)*delta2_1;
                weights(2,k) = weights(2,k) + coeff*input(j,k-1)*delta2_2;
                weights(3,k) = weights(3,k) + coeff*x2(k-1)*delta3_1;
            end
        end
    end
    mse(i) = (out - output)'*(out - output);
end
plot(mse)

```


Global minimum
run



Local minimum
run



Class + = ~~training~~ test error rate
 o = training error rate



- Note:
- 1) training error rate lower than test error, generally.
 - 2) Variance in performance grows with the number of hidden nodes
 - 3) Best test error occurs with ~ 24 hidden units,
 - 4) For this 6-class problem, test error rate is certainly much better than random guessing.